

# Troubleshooting Kubernetes applications

September 2021

# Contents

---

- 1. OVERVIEW ..... 1
- 2. TYPICAL APPLICATION STRUCTURE..... 1
- 3. HEALTH CHECK ON PODS ..... 2
- 4. SERVICE STATUS CHECK..... 12
- 5. INGRESS HEALTH CHECK ..... 14
- 6. QUICK INSPECTION GUIDE..... 16

# 1. Overview

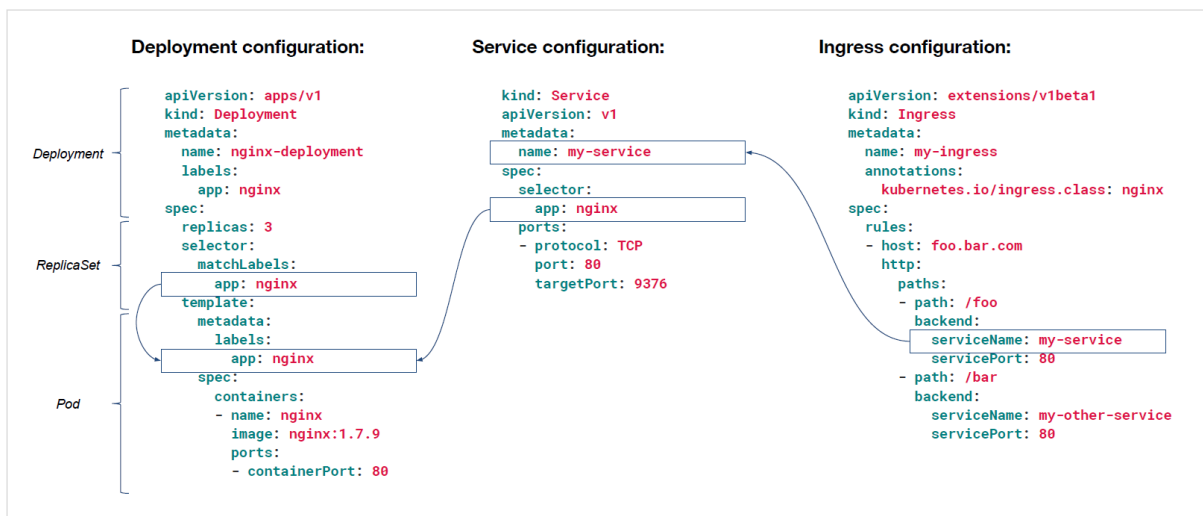
The purpose of this document is to describe how to troubleshoot issues while deploying workload on a Kubernetes cluster or operating the workload.

The most commonly used web application type will be introduced in the document. You can also refer to this document to troubleshoot SDS Cloud products of Kubernetes Engine and Kubernetes Apps, as well as vanilla Kubernetes and public managed Kubernetes including GKE, AKS, and EKS.

## 2. Typical application structure

Kubernetes offers various mechanisms such as hostNetwork, nodePort, and type: LoadBalancer to enable external access for web applications, but we will deal with the workload of the Ingress method, which is easy to access the http/https-based domain call method.

Here is an example Kubernetes resource combination for deploying an Nginx web application as a Kubernetes workload.



**Figure 1. Deployment strategies on Kubernetes, CNCF**

A variety of other resources such as persistent volume or network policy may exist but the example above is the most basic Kubernetes resource using Ingress.

The important thing here is the reference relationship between each resource. You must remember that both deployment and service selectors refer to the label of the pod template while serviceName refers to the name of the service resource in Ingress.

Next, we will look at normal status checks and responses to potential issues after the deployment of each resource.

## 3. Health check on pods

### 3.1 Basic commands

Here are some of the kubectl commands needed to check the status of a pod.

The following are the most basic query commands to check the basic pod status, an extended command including IP and assigned node information (-owide), or raw data in the form of a YAML manifest (-oyaml).

```
$ kubectl get pod
$ kubectl get pod -owide
$ kubectl get pod -oyaml
```

The following command is mainly used to check detailed information and events of pods if a problem occurs.

```
$ kubectl describe po <pod-name>
```

The command below can check the log if there is a problem with the running application. If multiple containers are running in the pod, you can select a specific container through the -c option, enable real-time tailing with the -f option and check logs for previous pods using the --previous option.

```
$ kubectl logs <pod-name> [-c <container-name>]
$ kubectl logs <pod-name> [-c <container-name>] -f
$ kubectl logs <pod-name> [-c <container-name>] --previous
```

The following is a command that shows the events that occurred in the namespace, which can be sorted by a chronological order with the --sort-by option.

```
$ kubectl get ev --sort-by=.metadata.creationTimestamp
```

### 3.2 Lifecycle of normal pods

After deploying a pod, the flow of the states until the pod changes to a healthy state has to be tracked. A typical pod has the following lifecycle.



- **Pending:** This status is displayed before determining which node will be scheduled to a pod or if there are no schedulable nodes.
- **ContainerCreating:** This stage is when the container runtime of a node creates

a container after a pod is scheduled on a specific node. It can be displayed in the process of mounting related resources, such as image pulling, ConfigMaps, and persistent volume.

- **Running:** The process in the pod is running normally.
- **Ready:** After the Running status, the pod is ready to provide service as per the readiness probe set in the pod.
- **Terminating:** This status may appear in the process of cleaning up a pod from the applicable node when the pod has been deleted.

You can check changes in the pod lifecycle by using the `-w` (watch) option in the `kubectl get po` command below. The most basic columns to check the normal status of the pod among these are `READY` and `STATUS`.

```
$ kubectl get po -owide -w
```

NAME	READY	STATUS	AGE	IP	NODE
nginx-5d796fc999-qkgpp	0/1	Pending	0s	<none>	<none>
nginx-5d796fc999-qkgpp	0/1	Pending	0s	<none>	node1
nginx-5d796fc999-qkgpp	0/1	ContainerCreating	1s	<none>	node1
nginx-5d796fc999-qkgpp	0/1	Running	3s	10.44.0.2	node1
nginx-5d796fc999-qkgpp	1/1	Running	9s	10.44.0.2	node1

Next, we will look at the causes of and responses for each case of abnormal pod status.

### 3.3 If Pending state persists

Sometimes a pod's state stays in Pending for a long time.

```
$ kubectl get po
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-6dcd8d4dff-njsrh	0/1	Pending	0	10m

In this case, Event messages need to be checked first using the `kubectl describe pod` command.

```
$ kubectl describe po <pod-name>
```

Events:

Type	Reason	Age	From
Warning	FailedScheduling	14s (x2 over 14s)	default-scheduler

0/3 nodes are available: **3 Insufficient cpu/memory.**

In the above case, the cluster is running out of resources. You should check the Allocated resources: assigned to each node through the `kubectl describe node` command, and if necessary, increase the cluster node resources or reduce the pod resources.

```
$ kubectl describe po <pod-name>
Events:
  Type      Reason           Age          From          Message
  ----      -
Warning    FailedScheduling  5s (x2 over 5s)  default-scheduler  0/3 nodes are available: 3
node(s) didn't match node selector.
```

In this case, there is no node matching the `nodeSelector/nodeAffinity` set in the pod. You can normally schedule a pod by modifying the `nodeSelector/nodeAffinity` of the pod to the matching node information or adding a label to the node.

```
$ kubectl describe po <pod-name>
Events:
  Type      Reason           Age          From          Message
  ----      -
Warning    FailedScheduling  6s (x2 over 6s)  default-scheduler  0/3 nodes are available: 3
node(s) were unschedulable.
```

In this case, the node is in an unschedulable state, as all nodes are cordoned or in `NotReady` state to cause the taint to be hanging. You must either uncorordon the node or act on the node in `NotReady` state.

```
$ kubectl describe po <pod-name>
Events:
  Type      Reason           Age          From          Message
  ----      -
Warning    FailedScheduling  53s (x2 over 60s)  default-scheduler  error while running
"VolumeBinding" filter plugin for pod "nginx-9d8d9896-k7gzj": pod has unbound immediate
PersistentVolumeClaims
```

You need to make sure that the PVC (`PersistentVolumeClaim`) is not in `Pending` state.

```
$ kubectl get pvc
NAME          STATUS    VOLUME   CAPACITY   ACCESS MODES   STORAGECLASS
nginx-pvc     Pending                                     standard
```

When using a static PV (`PersistentVolume`), the PV in the `Available` state is normally created first, and then you need to check whether the PV name properly matches to the `volumeName` of the PVC.



If you are using dynamic PV, you need to check if there is a StorageClass created with the storageClassName set in the PVC and then check if the Volume Provisioner is working properly.

```
$ kubectl describe po <pod-name>
Events:
  Type      Reason      Age   From              Message
  ----      -
  Normal    Scheduled   21s   default-scheduler Successfully assigned default/nginx-5d796fc999-6ksbj to node1
$ kubectl get po -owide
NAME                                READY   STATUS    RESTARTS   AGE   NODE   nginx-
5d796fc999-6ksbj                    0/1     Pending   0           37s   node1
```

For cases that only the event indicating a normal assignment of the pod to the node appears while the pod remains Pending, you need to check if the kubelet of the node that displays when you use `-owide` is working properly and take action accordingly.

```
$ kubectl describe po <pod-name>
Events: <none>
$ kubectl get po -n kube-system
NAME                                READY   STATUS              RESTARTS   AGE
kube-scheduler-master              0/1     CrashLoopBackOff    9           16m
```

If no message is recorded in the event, kube-scheduler is not properly scheduling and therefore measures need to be taken.

### 3.4 When ContainerCreating state persists

Below is what happens when a pod stays in the ContainerCreating state.

```
$ kubectl get po
NAME                                READY   STATUS              RESTARTS   AGE
nginx-6dcd8d4dff-njsrh             0/1     ContainerCreating   0           1m
```

Even in this case, you must first check the Event message using the `kubectl describe pod` command.

```
$ kubectl describe po <pod-name>
Events:
  Type      Reason      Age   From              Message
  ----      -
  Normal    Scheduled   <unknown>   default-scheduler Successfully assigned default/nginx-58cd5b85f7-cmb6c to node1
  Normal    Pulling     12s   kubelet, node1    Pulling image "nginx:1.13.9"
```

If the Event message is stuck at the pulling image state, the images are still being pulled and you need to wait a little longer. Pulling may take long if the image capacity is large or the network performance is low, and the issue will be resolved naturally when image pulling is completed.

```
$ kubectl describe po <pod-name>
Events:
  Type      Reason      Age           From          Message
  ----      -
Warning    FailedMount  5s (x6 over 21s)  kubelet, node2  MountVolume.SetUp failed for volume "config-volume" : configmap/secret "nginx-cm" not found
```

The case below is when the ConfigMap/Secret used by the pod is not created or an incorrect name is specified. You need to make sure that the ConfigMap/Secret specified in the volumes field of the pod spec actually exists.

```
$ kubectl describe po <pod-name>
Events:
  Type      Reason      Age           From          Message
  ----      -
Warning    FailedMount  18s          kubelet, node2  MountVolume.SetUp failed for volume "pvc-d2def29a-d503-443e-9615-886713983216" : mount failed: exit status 32
```

This occurs when a pod is not mounted to the PV in use. For NFS or Container Storage Interface (CSI) Driver to use the PV, you need to check if nfs-utils, rpcbind, or other clients are installed. In addition, the access control and firewall for the PV in the node need to be checked to enable the mount.

### 3.5 If ImagePullBackOff status is displayed

This time, we will look at the situation where ImagePullBackOff occurs.

```
$ kubectl get po
NAME                                READY  STATUS             RESTARTS  AGE
nginx-6dcd8d4dff-njsrh             0/1    ImagePullBackOff   0          6s
```

You can check it with the describe command.

```
$ kubectl describe po <pod-name>
Events:
  Type      Reason      Age           From          Message
  ----      -
Warning    Failed      23s (x2 over 44s)  kubelet, node1  Failed to pull image "nginx:invalid-tag": rpc error: code = Unknown desc = Error response from daemon: manifest for nginx:test not found
```



This is when there is no Image:Tag in the Image Repository. You need to edit the Image:Tag to the correct information in the pod specifications or push the Image:Tag to the Image Repository.

```
$ kubectl describe po <pod-name>
Events:
  Type      Reason      Age      From          Message
  ----      -
Warning    Failed      9s       kubelet, node1    Failed to pull image
"myregistry.io/nginx:1.17.8": rpc error: code = Unknown desc = Error response from daemon: Get
https://myregistry.io/v2/nginx/manifests/1.17.8: no basic auth credentials
```

If you are using a private registry, you may need credential information. In this case, you need to create imagePullSecrets with the correct credential information and add it to the pod spec. If imagePullSecrets is already in use, additional check on the expiration of authority is required.

```
$ kubectl describe po <pod-name>
Events:
  Type      Reason      Age      From          Message
  ----      -
Warning    Failed      12s     kubelet, node1    Failed to pull image "nginx:1.17.8": rpc error: code =
Unknown desc = Error response from daemon: Get https://registry-1.docker.io/v2/: net/http: request
canceled while waiting for connection (Client.Timeout exceeded while awaiting headers)
```

Above shows a failure of communications with the Image Repository in the node. You need to check if the firewall between Image Repositories and the Image Repository service are properly operating on the node.

### 3.6 If CrashLoopBackOff and Running repeat periodically

A pod may sometimes repeat the CrashLoopBackOff and Running states.

```
$ kubectl get po -w
NAME                                READY   STATUS             RESTARTS   AGE
test-7f74c45f58-jm629             0/1     CrashLoopBackOff   8          16m
test-7f74c45f58-jm629             0/1     Running            9          16m
```

If this takes place, you need to check the current or previous pod logs using the command below and troubleshoot the application if the log is confirmed.

```
$ kubectl logs po <pod-name> -f
$ kubectl logs po <pod-name> --previous
```

```
$ kubectl describe po <pod-name>
Last State:      Terminated
Reason:          ContainerCannotRun
Message:          OCI runtime create failed: container_linux.go:344: starting container process
caused "exec: \"ping\": executable file not found in $PATH": unknown
Exit Code:       127
```

The most common application issues include DB connection failure, abnormal command use, and application configuration issues.

If you cannot see the current or previous pod logs, check the Last State: details with the describe command.

### 3.7 If CrashLoopBackOff and Completed repeat periodically

Pods may also repeat CrashLoopBackOff and Completed states.

```
$ kubectl get po -w
```

NAME	READY	STATUS	RESTARTS	AGE
test-6fd77b68b9-44hxs	0/1	CrashLoopBackOff	8	16m
test-6fd77b68b9-44hxs	0/1	Completed	9	16m

In this case, there is no foreground process inside the container. You must specify a process that can operate in the foreground with CMD or ENTRYPOINT in the Dockerfile or add a command to Deployment. If it is a batch process, it must be created as a Job or CronJob.

### 3.8 If status does not change to Ready

The pod is running, but the status does not change to READY.

```
$ kubectl get po
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-659484b897-mnq4l	0/1	Running	0	2m14s

If this happens, use the describe command.

```
$ kubectl describe po <pod-name>
Events:
  Type      Reason      Age      From          Message
  ----      -
Warning    Unhealthy   5s (x7 over 65s)    kubelet, node2    Readiness probe failed: HTTP probe failed with statuscode: 404
```

The Event messages need to be checked to see if the readiness probe has failed, followed by necessary measures. If the application startup time is too long, `initailDelaySeconds` has to be increased.

### 3.9 If Terminating state continues

The next case is when the Terminating state persists.

```
$ kubectl get po -owide
```

NAME	READY	STATUS	RESTARTS	AGE	NODE
nginx-6dcd8d4dff-njsrh	0/1	Terminating	0	10m	node2
nginx-6dcd8d4dff-fvlmf	1/1	Running	0	43s	node1

First, we need to make sure that the node to which the pod is assigned is not in `NotReady` state. It is also important to note that in `NotReady` state, pods have different Eviction policies for different types of workloads.

```
$ kubectl get no
```

NAME	STATUS	ROLES	AGE	VERSION
master	Ready	master	39d	v1.21.2
node1	Ready	worker	39d	v1.21.2
node2	NotReady	worker	39d	v1.21.2

- Deployment: Check if it has failed over to another node
- StatefulSet: Not moved to another node (can be moved when forcibly deleted)
- DaemonSet: Not moved to another node

If the node to which the pod is assigned is `Ready`, the issue may be that PV is unmounted or detach does not work properly during the pod shutdown process, requiring a check on storage connection.

For other reasons, you should see if a zombie (defunct) process has been created on the node.

### 3.10 If Evicted

This is when the pod is Evicted.

```
$ kubectl get po -owide
```

NAME	READY	STATUS	RESTARTS	AGE	NODE
nginx-58cd5b85f7-4fz88	0/1	Evicted	0	1h	node1
nginx-58cd5b85f7-ts7ps	1/1	Running	0	12m	node2

You need to check the status by describing the node.

```
$ kubectl describe no node1
Taints:          node.kubernetes.io/disk-pressure:NoSchedule
Conditions:
  Type            Status  LastHeartbeatTime             LastTransitionTime
  Reason          Message
  ----            -
  DiskPressure    True    Tue, 20 Oct 2020 12:22:14 +0900  Wed, 14 Oct 2020 20:37:51
+0900    KubeletHasDiskPressure    kubelet has disk pressure
```

You should check if DiskPressure is caused by insufficient disk space on the node and if so, free up the space.

```
$ df -h /var/lib/kubelet
Filesystem              Size  Used Avail Use% Mounted on
/dev/mapper/VGROOT-LV_root  50G   46G   4.7G  91% /
```

Nodes are usually evicted when the disk to which the /var/lib/kubelet area belongs is 90% full or more (when using the default kubelet policy).

### 3.11 Restart frequently occurs

There are cases with frequent restarts even when the pod is running normally.

```
$ kubectl get po -n kube-system
NAME                READY  STATUS   RESTARTS  AGE
$ kubectl logs weave-net-kwzvz -n kube-system -c weave --previous
FATA: 2020/10/15 04:25:03.297881 [kube-peers] Could not get peers: Get
https://172.24.0.1:443/api/v1/nodes: dial tcp 172.24.0.1:443: i/o timeout
Failed to get peers
```

If the problem can be resolved by checking the previous pod logs with the --previous option, you should troubleshoot the application.

Common issues include communication failure with an external DB service or an internal pod related to the application causing restarts due to normal node work or network PM work. If the previous pod log cannot be confirmed, you need to check Last Stated: using the describe command.

```
$ kubectl describe po <pod-name>
Last State:      Terminated
Reason:         OOMKilled
Exit Code:      137
```

In the case of a shutdown due to OOM (OutOfMemory), pod resources should be configured considering service availability.

```
$ kubectl describe po <pod-name>
Events:
  Type       Reason      Age          From          Message
  ----       -
Warning      Unhealthy   48m (x6 over 49m)    kubelet, node2    Liveness probe failed: localhost:5432 - no response
```

If the liveness probe fails, the liveness probe threshold should be set in consideration of the load situation and the connection lease.

### 3.12 If pod itself cannot be queried

This is when you create a workload such as Deployment but the pod cannot be queried.

In this case, you need to look up the Event object first.

```
$ kubectl get ev
LAST SEEN   TYPE      REASON      OBJECT                                MESSAGE
3m1s        Warning   FailedCreate replicaset/nginx-847f85d779   Error creating: pods "nginx-847f85d779-j484f" is forbidden: exceeded quota: resourcequota, requested: requests.cpu=100m,requests.memory=2Gi, used: requests.cpu=4,requests.memory=128Mi, limited: requests.cpu=2,requests.memory=1Gi
```

In the case above, ResourceQuota is set in the Namespace and the corresponding pod exceeds the quota or the LimitRange. You need to adjust either the ResourceQuota/LimitRange or pod resources.

```
$ kubectl get ev
LAST SEEN   TYPE      REASON      OBJECT                                MESSAGE
3s          Warning   FailedCreate replicaset/nginx-56b5449445   Error creating: pods "nginx-56b5449445-" is forbidden: error looking up service account default/my-sa: serviceaccount "my-sa" not found
```

In this case, the resource required for running the pod has not been created. Necessary resources need to be created by looking up the pod specification.

```
$ kubectl get po -n kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
kube-controller-manager-master	0/1	CrashLoopBackOff	6	10m

If kube-controller-manager is abnormal, you need to check the log of kube-controller-manager and take necessary action.

## 4. Service status check

If there is no problem with the running state of the pod, you need to check the status of the service. First, you need to check if there is any issue in the communications to the pod.

### 4.1 Checking pod communication

Kubernetes nodes can directly send a communication request to <Pod IP>:<Port>.

```
$ kubectl get po -owide
```

NAME	READY	STATUS	RESTARTS	AGE	IP
<pre>\$ kubectl port-forward &lt;pod-name&gt; &lt;local-port&gt;:&lt;pod-port&gt;</pre> <pre>Forwarding from 127.0.0.1:8888 -&gt; 80</pre> <pre>Forwarding from [::1]:8888 -&gt; 80</pre> <pre>\$ curl localhost:8888</pre> <pre>&lt;title&gt;Welcome to nginx!&lt;/title&gt;</pre>					

If the node is not a Kubernetes node, the communication status may be checked through port-forward of kubectl as shown below (only for TCP).

After pod communication is confirmed, you can check service communication.

### 4.2 Pod communication failure

However, if communication with the pod is not available, you need to first check if the pod is listening to the corresponding port.

In the case of normal port listening, the cause may be a communication block by network policies. You need to check communication from other pods that have accepted network policies.

```
$ kubectl get netpol
```

NAME	POD-SELECTOR	AGE
default-deny	<none>	1d

Alternatively, you could check if the Container Network Interface (CNI) Plugin is operating normally and take appropriate action.

```
$ kubectl get po -n kube-system | grep weave
```

NAME	READY	STATUS	RESTARTS	AGE
weave-net-97gcr	1/2	CrashLoopBackOff	8	16m

### 4.3 Check service communication

As with pods, Kubernetes nodes also allow direct communication checks with <Service IP>:<Port>.

```
$ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
nginx	ClusterIP	10.100.215.120	<none>	80/TCP	54m

```
$ curl 10.100.215.120:80
```

<title>Welcome to nginx!</title>

If they are not Kubernetes nodes, you can check communication status through kubectl port-forward (only for TCP).

```
$ kubectl port-forward service/<service-name> <local-port>:<pod-port>
```

Forwarding from 127.0.0.1:8888 -> 80  
Forwarding from [::1]:8888 -> 80

```
$ curl localhost:8888
```

<title>Welcome to nginx!</title>

If service communication has been confirmed up to this point, you can proceed to Ingress communication check.

### 4.4 Service communication failure

If communication with the service fails, you can try the following inspection items.

```
$ kubectl describe svc <service-name>
```

Selector: name=nginx  
Port: http 80/TCP  
TargetPort: 8080/TCP  
Endpoints: <none>

When describing a service, it is necessary to check whether or not endpoints have been used. If it is <none>, the service selector and the pod label may have been mismatched and they must be corrected.



Below displays a communication failure even with an endpoint.

```
$ kubectl describe svc <service-name>
Selector:          name=nginx
Port:             http 80/TCP
TargetPort:       8080/TCP
Endpoints:        10.36.0.1:8080
```

If the targetPort of the service and the containerPort of the pod are mismatched, the ports must be modified to match. If this is not the case, you should check if communication is disabled due to network policies, see if kube-proxy is abnormal, and then take following measures.

## 5. Ingress health check

If maintenance on the service has been completed, you can finish checking workloads related to Kubernetes after checking communication with Ingress.

### 5.1 Check Ingress communication

Under normal circumstances, pod IP information should be present in the backend during describing.

```
$ kubectl describe ing <ingress-name>
Name:          nginx
Address:       172.28.128.11
Rules:
  Host          Path  Backends
  ----          -
  nginx.example.io
                  nginx:80 (10.36.0.1:80,10.44.0.2:80)
```

You can check the domain set in Ingress by directly accessing it through a web browser or with the curl command on the server. If the domain is not registered in the DNS server, you can check it by adding the domain to the header as shown below.

```
$ curl -H 'Host: <host-domain>' http://<ingress-controller-external-ip>
<title>Welcome to nginx!</title>
```

If communications have worked up to this point, there seems to be no issue within the Kubernetes cluster.

### 5.2 Ingress communication failure

There are also cases where communication to Ingress fails.

First, there could be no pod IP information in the Ingress backend.

```
$ kubectl describe ing <ingress-name>
Name:          nginx
Address:       172.28.128.11
Rules:
  Host          Path  Backends
  ----          -
  nginx.example.io
                  nginx:80 (<none>)
```

In this case, you should confirm if the serviceName and servicePort in the backend of the Ingress spec are properly matched with the actual service resource.

From here on are cases where communication with Ingress backend information is not working.

```
$ curl -H 'Host: nginx.example.io' http://172.28.128.11
<title>504 Gateway Time-out</title>
```

For the case below, you need to check if communication has been blocked by network

```
$ curl -H 'Host: nginx.example.io' http://172.28.128.11
<title>503 Service Temporarily Unavailable</title>
```

policies.

If the application service is not responding, you need to check if the pod status is

```
$ curl -H 'Host: nginx.example.io' http://172.28.128.11
<title>404 Not Found</title>
```

normal.

You must check if the path set in the Ingress backend is the context path provided by the actual application.

```
$ curl -H 'Host: nginx.example.io' http://172.28.128.11
<title>404 Not Found</title>
```

You need to make sure that the Ingress Controller and firewall are open.

```
$ curl -H 'Host: nginx.example.io' http://172.28.128.11
curl: (7) Failed connect to 172.28.128.11:80; Connection refused
```

This may occur when the Ingress Controller does not listen to the port. You should check if the Ingress Controller is running normally.

Additionally, if it is blocked by a firewall from the outside, check communication through the command below and then open <ingress-controller-external-ip>:80,443 and the firewall.

```
$ telnet <ingress-controller-external-ip> 80
or
$ </dev/tcp/<ingress-controller-external-ip>/80
```

If the domain is not registered externally, a "Could not resolve host" error may occur. After checking if the domain issue can be resolved using the following command, you need to register the domain in the DNS server or individual domains in the host files before use.

```
$ nslookup <host-domain>
or
$ getent hosts <host-domain>
```

## 6. Quick inspection guide

The following is the order of a quick inspection based on the guidelines so far.

Pod Running

Pod Ready

Pod Conn

Service Conn

Ingress Conn

### 1. Check Running and Ready states of pods

```
$ kubectl get po
NAME                                READY   STATUS    RESTARTS   AGE
nginx-58cd5b85f7-84xlr             1/1     Running   0           31s
```

### 2. Check pod communication

```
$ curl <Pod IP>:<port>
```

### 3. Check service communication

```
$ curl <Service IP>:<port>
```

### 4. Check Ingress communication

```
$ curl <Service IP>:<port>
```